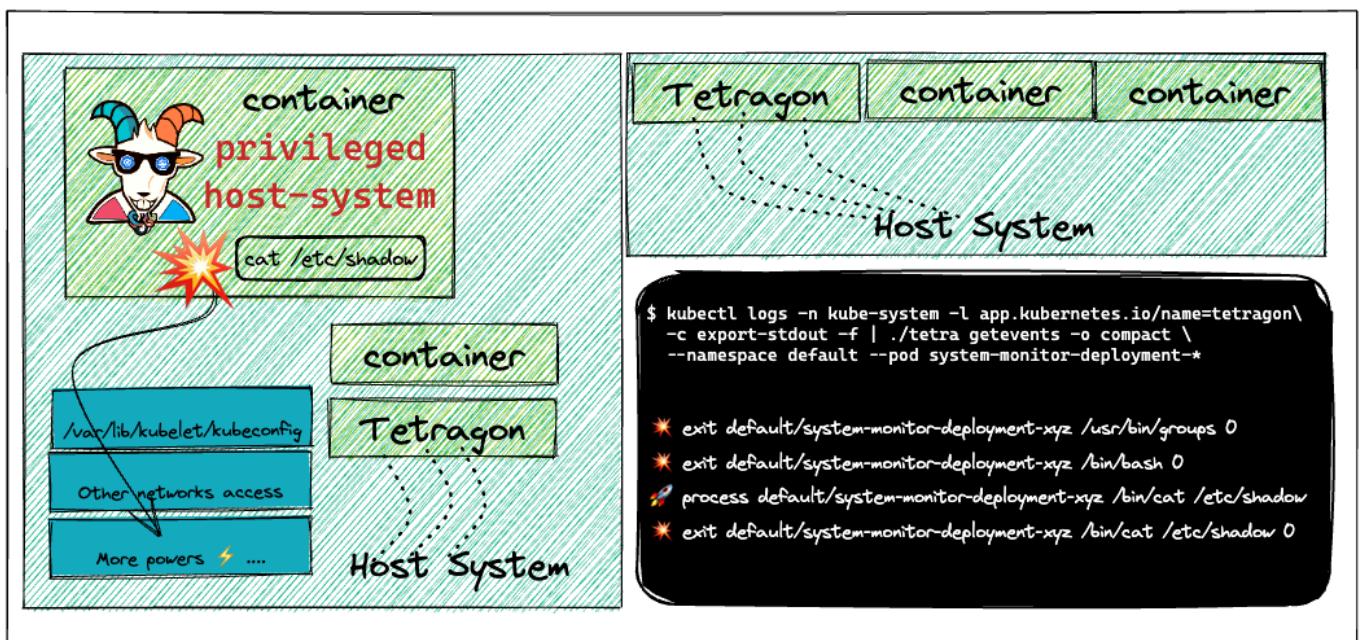


✳️ Cilium Tetragon - eBPF-based Security Observability and Runtime Enforcement

🙌 Overview

The containers and their infrastructure are immutable. It means it's very difficult to detect certain attacks, vulnerabilities, and detections using traditional tools and technologies. In this scenario, we will see how we can leverage the popular open-source tool like Cilium Tetragon to detect and perform runtime security monitoring using the `tracingpolicy` in action.



By the end of the scenario, we will understand and learn the following

1. Deploying the helm chart into the Kubernetes cluster
2. Performing the log analysis and detection of security events in the Kubernetes Cluster
3. Use, analyze and detect security issues in near real-time using Cilium Tetragon

⚡ The story

This scenario is to deploy runtime security monitoring & detection for containers and Kubernetes resources. Also, explore and see how we can detect certain issues and perform detections using `tracingpolicy` based on attacker triggers.

NOTE

Make sure you run the following deployment using Helm with v3. Refer to [helm installation](#)

INFO

- To deploy the Cilium Tetragon helm chart run the following commands

```
helm repo add cilium https://helm.cilium.io
```

```
helm repo update
```

```
helm install tetragon cilium/tetragon -n kube-system
```

- You can verify that the Tetragon pods are in running state by running the following command

```
kubectl rollout status -n kube-system ds/tetragon -w
```

```
$ helm repo add cilium https://helm.cilium.io
"cilium" has been added to your repositories
$ helm repo update
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "cilium" chart repository
Update Complete. Happy Helming!
$ helm install tetragon cilium/tetragon -n kube-system
NAME: tetragon
LAST DEPLOYED: Sat Apr 8 16:27:40 2023
NAMESPACE: kube-system
STATUS: deployed
REVISION: 1
TEST SUITE: None
$
$ kubectl rollout status -n kube-system ds/tetragon -w
daemon set "tetragon" successfully rolled out
$
```

🎯 Goal



TIP

Deploy the Cilium Tetragon and detect security events in near real-time using the `tracingpolicy` and events.

◻ Hints & Spoilers

- ▶ ✨ Want to know more?

🎉 Solution & Walkthrough

🎲 Method 1



INFO

Cilium's new Tetragon component enables powerful realtime, eBPF-based Security Observability and Runtime Enforcement.

Tetragon detects and is able to react to security-significant events, such as

- Process execution events
- System call activity
- I/O activity including network & file access

When used in a Kubernetes environment, Tetragon is Kubernetes-aware - that is, it understands Kubernetes identities such as namespaces, pods and so-on - so that security event detection can be configured in relation to individual workloads.

- Get more details about the Tetragon deployment by running the following command

```
kubectl get pods -n kube-system --selector app.kubernetes.io/instance=tetragon
```

- Manually obtaining the logs from the Tetragon using the following command

```
kubectl logs -n kube-system -l app.kubernetes.io/name=tetragon -c export-stdout -f
```

- Now, let's use the system-monitor pod and try privileges escalation to the host system and see if that detects by Tetragon

```
export POD_NAME=$(kubectl get pods -l "app=system-monitor" -o jsonpath=".items[0].metadata.name")
kubectl exec -it $POD_NAME bash
```

- Let's read the sensitive file `/etc/shadow`

```
cat /etc/shadow
```

```
root-plane", "time": "2023-04-08T16:50:09.320157831Z"}
{"process_exit": {"process": {"exec_id": "a3ViZXJuZXRLcy1nb2F0LWNvbnRyb2wtcGxhbmU6MTUyMTE0NjY30DgxNTM6MTI0MjA=", "pid": 12420, "uid": 0, "cwd": "/", "binary": "/bin/cat", "arguments": ["/etc/shadow"], "flags": "execute_rroot cwd clone", "start_time": "2023-04-08T16:50:09.320057942Z", "audit": 4294967295, "pod": {"namespace": "default", "name": "system-monitor-deployment-674bb4dc65-wtq5q"}, "container": {"id": "containererd://401efe141e0932e1e756314a315d90a3e6c0a31f7e6e48cab0a8e1dd8508310a"}, "name": "system-monitor", "image": {"id": "docker.io/madhuakula/k8s-goat-system-monitor@sha256:06b58bd080201ea0d4048befdd2159f384b61ce457a5a96e3001db629b5caa40"}, "name": "docker.io/madhuakula/k8s-goat-system-monitor:latest"}, "start_time": "2023-04-08T16:38:01Z", "pid": 5371}, "pod_labels": {"app": "system-monitor", "pod_template_hash": "674bb4dc65"}, "docker": "401efe141e0932e1e756314a315d90a", "parent_exec_id": "a3ViZXJuZXRLcy1nb2F0LWNvbnRyb2wtcGxhbmU6MTUxNjg4MDDE1MjY20Dc6MTIyOTc=", "parent": {"exec_id": "a3ViZXJuZXRLcy1nb2F0LWNvbnRyb2wtcGxhbmU6MTUxNjg4MDDE1MjY20Dc6MTIyOTc=", "pid": 12297, "uid": 0, "cwd": "/", "binary": "/bin/bash", "flags": "execute_rroot cwd clone", "start_time": "2023-04-08T16:49:26.654797498Z", "audit": 4294967295, "pod": {"namespace": "default", "name": "system-monitor-deployment-674bb4dc65-wtq5q"}, "container": {"id": "containererd://401efe141e0932e1e756314a315d90a3e6c0a31f7e6e48cab0a8e1dd8508310a"}, "name": "system-monitor", "image": {"id": "docker.io/madhuakula/k8s-goat-system-monitor@sha256:06b58bd080201ea0d4048befdd2159f384b61ce457a5a96e3001db629b5caa40"}, "name": "docker.io/madhuakula/k8s-goat-system-monitor:latest"}, "start_time": "2023-04-08T16:38:01Z", "pid": 5258}, "pod_labels": {"app": "system-monitor", "pod_template_hash": "674bb4dc65"}, "docker": "401efe141e0932e1e756314a315d90a", "parent_exec_id": "a3ViZXJuZXRLcy1nb2F0LWNvbnRyb2wtcGxhbmU6MTUxNjg6MTIyODQ=", "refcnt": 1, "time": "2023-04-08T16:50:09.34026625Z"}, "node_name": "kubernetes-goat-control-plane", "time": "2023-04-08T16:50:09.340265784Z"}
```

```
root@system-monitor-deployment-674bb4dc65-wtq5q:/# cat /etc/shadow
```

- You can also see these events in a nicer way using the official `tetra` cli client in your local system. Refer to the documentation for specific binary as per your **OS and Architecture**

```
wget https://github.com/cilium/tetragon/releases/download/v0.9.0/tetra-linux-amd64.tar.gz
tar -xvzf tetra-linux-amd64.tar.gz
```

- Now you can run the following command to pass the output of the Tetragon events to `tetra` cli locally to see in a nicer way

```
kubectl logs -n kube-system -l app.kubernetes.io/name=tetragon -c export-stdout -f | ./tetra getevents -o compact --namespace default --pod system-monitor-deployment-*
```

The screenshot shows a terminal window with two main sections. The top section displays Tetragon event logs for the 'system-monitor' deployment. The bottom section shows the contents of the '/etc/shadow' file on the host system.

```
root@system-monitor-deployment-674bb4dc65-wtq5q:/# ./tetra getevents -o compact --namespace default --pod system-monitor-deployment-*
* exit  default/system-monitor-deployment-674bb4dc65-wtq5q /usr/bin/groups  0
* process default/system-monitor-deployment-674bb4dc65-wtq5q /usr/bin/dircolors -b
* exit  default/system-monitor-deployment-674bb4dc65-wtq5q /usr/bin/dircolors -b 0
* exit  default/system-monitor-deployment-674bb4dc65-wtq5q /bin/bash  0
* process default/system-monitor-deployment-674bb4dc65-wtq5q /bin/cat /etc/shadow
* exit  default/system-monitor-deployment-674bb4dc65-wtq5q /bin/cat /etc/shadow  0

root@system-monitor-deployment-674bb4dc65-wtq5q:/# cat /etc/shadow
root:*:19110:0:99999:7:::
daemon:*:19110:0:99999:7:::
bin:*:19110:0:99999:7:::
sys:*:19110:0:99999:7:::
sync:*:19110:0:99999:7:::
games:*:19110:0:99999:7:::
man:*:19110:0:99999:7:::
lp:*:19110:0:99999:7:::
mail:*:19110:0:99999:7:::
news:*:19110:0:99999:7:::
uucp:*:19110:0:99999:7:::
proxy:*:19110:0:99999:7:::
```

TIP

`tetra` cli provides the context awareness of the Kubernetes, namespaces and other details like processes, etc. When querying you can even limit them to processes, namespace, pod, even regex supported.

- Let's take a spin for detecting the privilege escalation attacks using the Tetragon. You can perform the container escape using the `system-monitor` pod to gain host system access by running the following commands

```
# Get into the system-monitor pod
export POD_NAME=$(kubectl get pods -l "app=system-monitor" -o jsonpath=".items[0].metadata.name")
kubectl exec -it $POD_NAME bash
```

```
# Exploit to gain the host system access using nsenter  
nsenter -t 1 -m --uts --ipc --pid
```

```
root@system-monitor-deployment-59b94676c7-db9vl:/# hostname  
system-monitor-deployment-59b94676c7-db9vl  
root@system-monitor-deployment-59b94676c7-db9vl:/# nsenter -t 1 -m --uts --ipc --pid  
# hostname  
kubernetes-goat-control-plane  
# █
```

- As you can see following, Tetragon in action detecting these attacks in near real-time

```
kubectl logs -n kube-system -l app.kubernetes.io/name=tetragon -c export-  
stdout -f | ./tetra getevents -o compact --namespace default --pod system-  
monitor-deployment-*
```

```
process default/system-monitor-deployment-59b94676c7-db9vl /binhostname  
open default/system-monitor-deployment-59b94676c7-db9vl /binhostname /etcld.so.cache  
close default/system-monitor-deployment-59b94676c7-db9vl /binhostname  
exit default/system-monitor-deployment-59b94676c7-db9vl /binhostname 0  
process default/system-monitor-deployment-59b94676c7-db9vl /usrbinnsenter -t 1 -m --uts --ipc --pid  
open default/system-monitor-deployment-59b94676c7-db9vl /usrbinnsenter /etcld.so.cache  
close default/system-monitor-deployment-59b94676c7-db9vl /usrbinnsenter  
process default/system-monitor-deployment-59b94676c7-db9vl /binsh  
open default/system-monitor-deployment-59b94676c7-db9vl /binsh /etcld.so.cache  
close default/system-monitor-deployment-59b94676c7-db9vl /binsh  
open default/system-monitor-deployment-59b94676c7-db9vl /binsh /etcprofile  
close default/system-monitor-deployment-59b94676c7-db9vl /binsh  
process default/system-monitor-deployment-59b94676c7-db9vl /usrbinid -u  
open default/system-monitor-deployment-59b94676c7-db9vl /usrbinid /etcld.so.cache  
close default/system-monitor-deployment-59b94676c7-db9vl /usrbinid  
open default/system-monitor-deployment-59b94676c7-db9vl /binsh /etcprofile.d  
close default/system-monitor-deployment-59b94676c7-db9vl /binsh  
open default/system-monitor-deployment-59b94676c7-db9vl /binsh /etcprofile.d01localefix.sh  
close default/system-monitor-deployment-59b94676c7-db9vl /binsh  
exit default/system-monitor-deployment-59b94676c7-db9vl /usrbinid -u 0  
process default/system-monitor-deployment-59b94676c7-db9vl /usrbinlocalecheck C.UTF-8  
open default/system-monitor-deployment-59b94676c7-db9vl /usrbinlocalecheck /etcld.so.cache  
close default/system-monitor-deployment-59b94676c7-db9vl /usrbinlocalecheck  
process default/system-monitor-deployment-59b94676c7-db9vl /usrbinmsg n  
open default/system-monitor-deployment-59b94676c7-db9vl /usrbinmsg /etcld.so.cache  
close default/system-monitor-deployment-59b94676c7-db9vl /usrbinmsg  
exit default/system-monitor-deployment-59b94676c7-db9vl /usrbinlocalecheck C.UTF-8 0  
exit default/system-monitor-deployment-59b94676c7-db9vl /usrbinmsg n 1  
process default/system-monitor-deployment-59b94676c7-db9vl /usrbinhostname  
open default/system-monitor-deployment-59b94676c7-db9vl /usrbinhostname /etcld.so.cache  
close default/system-monitor-deployment-59b94676c7-db9vl /usrbinhostname  
exit default/system-monitor-deployment-59b94676c7-db9vl /usrbinhostname 0
```

- Hooray 🎉, now we can see that Tetragon detected the attack and notified using the events

References

- <https://github.com/cilium/tetragon>
- <https://tetragon.cilium.io>

